

Flow control makes shared inference predictable.

Flow control is the Endpoint Picker's policy layer for multi-tenant inference. While the pool has capacity, requests move through without being held. When the pool is under pressure, flow control activates. New work waits in policy-aware queues in front of the model servers, so priority, fairness, and tenant rules decide what runs next.

This page explains the mechanism end to end. What breaks without it, how a request travels through it, what it guarantees, what it costs, and how to operate it.

Why LLM traffic breaks rate limiting

Traditional gateways limit requests per second and assume every request costs about the same. LLM requests break that assumption.

Cost is driven by input length and by the autoregressive decode loop. One request can consume orders of magnitude more compute and memory than another at the same request rate.

Same request rate, very different cost

256-token prompt, 64-token reply

occupies a batch slot briefly

2048-token prompt, 512-token reply

occupies compute and KV cache for the whole decode loop

Measured on one GPU at 160 concurrent requests, the short shape cleared 165 requests per second and the long shape cleared 27. A request-per-second limit cannot see this difference. Physical capacity can.

Noisy neighbor

One tenant with heavy prompts monopolizes KV cache and queue space, starving the rest.

Scheduling regret

Once a request enters a server's local queue it cannot move, even when a better endpoint frees up.

Priority inversion

Model servers batch by arrival, not business priority. Real-time work gets stuck behind offline batches.

Resource asymmetry

One huge-context request can outweigh hundreds of small ones. Request counts miss it.

Flow control answers by governing physical capacity, KV cache and queue saturation, instead of request counts, and by holding the queue where policy can still act on it.

What flow control is, and what it is not

It is

A centralized, policy-aware queue in the Endpoint Picker, in front of the model servers. It enforces priority across

It is not

More capacity. Flow control decides when and in what order requests dispatch, and it cannot make the pool

traffic classes, fairness across tenants, and ordering within a tenant.

It is work-conserving. If the pool has capacity, it dispatches.

It routes late. A dispatched request lands on the endpoint that just freed up or holds its prefix cache.


faster.

When demand exceeds capacity, the wait still exists. Flow control moves it into the EPP queue and decides where and for whom time to first token is accrued, protecting generation speed for everything already running.

That second sentence is the honest core of the system. Enabling flow control is an explicit choice to protect time per output token, the reading pace of every stream already in flight, at the price of queue time for new arrivals when the pool is full.

The journey of a request

Five stops between the client and the GPU. The first two classify, the middle two decide, the last one routes.



The three-tier decision

Each dispatch cycle answers three questions in strict order. One is fixed, two are pluggable.

TIER	QUESTION	WHO DECIDES	OPTIONS
1 • Priority	Which band goes first?	Fixed. Highest band with work always wins	Bands come from <code>InferenceObjective</code> , negative values mark sheddable work
2 • Fairness	Which tenant inside the band?	Fairness policy plugin	Round-robin across tenants, or global strict ordering
3 • Ordering	Which request from that tenant?	Ordering policy plugin	First-come first-served, earliest deadline first, or SLO-deadline

Priority is absolute by design. When the pool is saturated, every premium request dispatches before any standard request, and fairness only referees between equals.

If two workloads should compete, put them in the same band. If one should never wait for the other, put them in different bands.

Where the waiting goes

When demand exceeds capacity, the wait has to land somewhere. There are only two places, and only one of them obeys policy.

Two waits, one budget

THE EPP QUEUE · policy decides

Time here is assigned. Priority picks the band, fairness picks the tenant, and the tenant that caused the pressure carries the most of it. Measurable per tenant and per band.

`llm_d_epp_flow_control_request_queue_duration_seconds`

THE RUNNING BATCH · physics decides

Once dispatched, a request computes inside the shared batch. A full batch produces first tokens more slowly for every request in it, whoever sent them. No scheduler reaches inside.

`vllm:num_requests_running`, `vllm:num_requests_waiting`

Tuning the saturation gate moves the wait between these two places. That is the dial. Let the batch fill for maximum utilization, or hold work upstream for maximum isolation, and choose per pool.

What the policy looks like in data

You can see each tier working from four measurements on one GPU running GPT-OSS 20B behind the llm-d gateway.

PRIORITY

41 vs 379 ms

p50 TTFT, premium vs standard

Standard surged past saturation. Premium stayed ahead of the queue it did not cause.

FAIRNESS

within 9 ms

equal peers, whole run

Two tenants at equal load were indistinguishable for 300 s while a third spiked.

COST ASSIGNMENT

107 vs 80 ms

mean queue time, spiker vs peers

The tenant that created the pressure carried the most queue time, every repeat.

SHEDDABLE ISOLATION

202 vs 64 ms

mean queue time, batch vs premium

Batch at triple the interactive arrival rate absorbed the waiting instead of causing it.

The absolute numbers track that specific GPU and tuning. The pattern is the guarantee.

Targets hold while capacity lasts. Past capacity the policy, not arrival order, decides who waits.

When requests are refused

Flow control fails loudly and specifically. Every refusal carries a reason on the wire in the `x-llm-d-request-dropped-reason` header.

WHAT HAPPENED	STATUS	REASON ON THE WIRE
A band's queue limit was hit, controlled load shedding	429	<code>rejected-saturated</code>

WHAT HAPPENED	STATUS	REASON ON THE WIRE
A request waited past its TTL	503	<code>rejected-ttl-expired</code>
The client disconnected while queued	503	<code>rejected-context-cancelled</code>
Graceful shutdown drained the queue, retryable	503	<code>rejected-shutting-down</code>

Queues are in-memory. If the EPP itself dies, the default llm-d configuration fails open, Envoy routes directly to the servers, and requests keep flowing without protection until the EPP returns.

Operating it

The knobs that matter

Band priorities in `InferenceObjective` map business tiers to dispatch order. **Per-band limits** bound how much low-priority work can queue.

Request TTL caps how stale queued work can get.

The saturation detector's threshold sets the healthy buffer, deep enough that batching never starves, shallow enough that the queue stays where policy lives. It is a trigger, not a queue cap.

What to watch

`pool_saturation` tells you when the gate is active.

`queue_size` and `queue_duration` per tenant and band tell you who is waiting and for how long. The same depth is a true-demand signal an autoscaler can act on.

On the server side, `num_requests_running` against `num_requests_waiting` tells you where the pressure actually sits.

One sweep per hardware and model pair finds the knee, the point where queueing begins. Size latency-sensitive pools below it, and let flow control govern everything that arrives past it.

That is the whole operating model. Know your envelope, place your operating point, and let the policy own the overflow.